

EASTERN AFRICA STATISTICAL TRAINING CENTRE



BACHELOR OF DATA SCIENCE YEAR III

PROJECT PROPOSAL

Title:

“Developing a Predictive Deep Learning Model for the Early Detection of Maize Leaf Diseases in Tanzania”

BY

1. SEVERIA M. BANDA	EASTC/BDTS/24/01380
2. MOSHI M. KASSIM	EASTC/BDTS/23/00309
1. GIDION F. MACHUMU	EASTC/BDTS/23/00440
2. REBECA P. SHAURI	EASTC/BDTS/24/01587
3. EDGAR S. MATERU	EASTC/BDTS/24/01607

Academic Year: 2025/2026

Supervisor's Name: Mr Kelvin Rushobola

Signature: _____

Submission Date: 16/06/2026

CHAPTER 1: INTRODUCTION

1.1 Background of the Study

Agriculture remains the backbone of the Tanzanian economy, employing approximately 55% of the national workforce (NBS & OCGS, 2024) and serving as the primary driver of national food security. The demand for food is growing exponentially and the production of agronomy is far too low. To overcome this, farmers, scientists, researchers, analysts, specialists, and the government try to place further effort and strategies to increase agricultural production to accommodate the needs. One of the major causes of low yield is diseases, which reduce the quality, quantity, and nutritional value of fruits, vegetables, cereals, and legumes, nearly 32% of the losses are observed due to diseases in cereal crops (Savary et al., 2019). At the heart of this agricultural sector is maize (*Zea mays*), the most widely cultivated and consumed staple crop across both rural and urban demographics. Maize is very essential crop all around the world, it is used for different purposes i.e., poultry, cattle feed, food, beverages, and it is prone to various diseases which results in low production (Savary et al., 2019). However, maximizing maize production is much hindered by the prevalence of destructive foliar diseases, most notably Common Rust, Northern Leaf Blight, Cercospora Leaf Spot, Maize Streak Virus and Maize Lethal Necrosis (Nsibo et al., 2024). These Phyto pathological threats can devastate entire harvests if not detected and managed during their early onset stages, with the potential to reduce yield.

The detection and classification of these maize diseases relies heavily on the manual scouting of fields and the visual expertise of agricultural extension officers. This traditional diagnostic pipeline is fundamentally constrained in a modern and scalable context, these methods have been observed to identify diseases at a later stage, potentially leading to more harm to the crops. Tanzania faces a shortage of localized agronomic experts, resulting in an imbalanced ratio of extension workers to smallholder farmers. By the time a farmer identifies anomalous lesions on a crop leaf, requests an expert evaluation and receives a definitive diagnosis, the pathogen has often spread beyond the point of localized intervention. This reliance on slow, subjective human visual inspection results in catastrophic delays in applying targeted phytosanitary treatments. Recently, technology has been used to improve yields in agriculture, whereby researchers have devised several solutions, including image processing and object detection using deep learning models (Mohanty et al., 2016). The rapid evolution of Artificial Intelligence and applied Data Science has introduced a paradigm shift in how biological data is processed.

Specifically, Computer Vision which is powered by Convolutional Neural Networks (CNNs) has proven unique capabilities of extracting complex and multi dimensional patterns from unstructured digital images(Mahlein, 2016). In the agricultural domain, deep learning algorithms can now be trained on thousands of labelled plant images to autonomously recognize microscopic and macroscopic disease manifestations with superhuman accuracy (Khan et al., 2023). By bridging the gap between advanced computational modelling and agronomy, data science offers a profound opportunity to digitize, automate and democratize crop diagnostics. Transforming a standard digital image into a highly accurate, real-time predictive output that has the potential to eliminate the traditional diagnostic bottleneck, ultimately securing maize yields against preventable biological shocks.

1.2 Problem Statement

Despite the critical importance of maize to Tanzania's food security, smallholder farmers face disproportionate yield losses due to the rapid and often undetected spread of foliar diseases such as Northern Leaf Blight and Common Rust. The fundamental problem lies in the structural bottleneck of traditional agricultural diagnostics. Currently, disease identification relies almost entirely on human visual inspection. Farmers, who may lack specialized Phyto pathological training, must wait for agricultural extension officers to physically visit their geographical isolated plots to provide an authoritative diagnosis. Tanzania suffers from a severe deficit of extension workers. This logistical constraint means that timely, expert and on-site evaluations are much rare during the crucial early stages of an outbreak. So, farmers are forced into a reactive rather than proactive agricultural stance. Without rapid and accurate disease identification, they just resort to misapplying broad-spectrum pesticides based on guesswork. This trial-and-error approach is not only financially crippling for low-income rural households and environmental degrading, but it also allows the specific underlying pathogen to spread unchecked, ultimately decimating the harvest. Laboratory based diagnostics, while scientifically accurate, are centralized in urban research institutions, rendering them inaccessible and far too slow for the average rural farmer who requires immediate intervention.

1.3 General Objective

To develop a Deep Learning image classification model capable of automatically identifying and diagnosing common maize leaf diseases which are Common Rust, Northern Leaf Blight and Cercospora Leaf Spot.

1.4 Specific Objectives

- To source a standardized dataset of healthy and diseased maize leaf images.
- To train a Convolutional Neural Network (CNN) model using transfer learning for classifying distinct maize foliar diseases.
- To evaluate the algorithm's predictive performance using standard data science metrics (Accuracy, Precision, Recall, and Confusion Matrix).
- To deploy the trained prediction model into an interactive graphical user interface (dashboard) for realtime image diagnosis

1.5 Research Questions

- What image preprocessing techniques are required to optimize raw leaf data for deep learning ingestion?
- How accurately can a fine tuned Convolutional Neural Network classify different strains of maize leaf diseases?
- How can the predictive model be effectively deployed to accept new, unseen images and output real-time diagnostic confidence scores?

1.6 Significance of the Study

This study aims to contribute in collective effort for protecting global food security and enhance agricultural productivity by modernizing how maize diseases are detected and managed, by overcoming the limitations of traditional methods, enabling early and precise detection of physiological changes in plants before symptoms are even visible to the naked eye.

CHAPTER 2: LITERATURE REVIEW.

2.1 Theoretical Framework

The computational diagnosis of agricultural diseases relies on a paradigm shift from traditional rule-based programming to data-driven learning. This study is underpinned by two foundational theories in Artificial Intelligence that explain how machines can be trained to recognize biological patterns. In traditional machine learning, human engineers must manually extract and define features e.g, writing code to identify the color brown or the geometric shape

of a leaf lesion. However, biological manifestations of crop diseases are highly variable and complex. To overcome this, this study relies on advanced cognitive computational theories that allow algorithms to autonomously learn these complexities directly from raw visual data(Wu, 2017).

Connectionism (Artificial Neural Networks): Connectionism is a cognitive science theory that models mental or behavioral phenomena as the emergent processes of interconnected networks of simple units(Medler, 1998). In computer science, this is operationalized as Artificial Neural Networks (ANNs). ANNs consist of layers of artificial "neurons" (nodes) that receive input, apply mathematical weights, and pass the signal to the next layer. This theory forms the base architecture of the diagnostic model. The algorithm will process maize leaf images by adjusting the mathematical weights across its network during training until it successfully links the visual input (pixels) to the correct output (disease diagnosis).

Representation Learning (via Convolutional Neural Networks): Representation Learning theory posits that a system can automatically discover the representations or features needed for detection or classification from raw data(Noroozi et al., 2017). Convolutional Neural Networks (CNNs) are the mathematical application of this theory for grid-like data (images). CNNs use convolutional filters to autonomously learn spatial hierarchies of features starting with simple edges and building up to complex disease patterns. This theory justifies the use of raw image data. Instead of telling the computer what "Northern Leaf Blight" looks like, the CNN will use representation learning to autonomously discover the unique visual signatures of the disease from the Plant dataset.

Relevance and Limitations of the Theories: The primary relevance of these theories is their unmatched accuracy in visual classification tasks, making them perfectly suited for identifying foliar plant diseases. However, a major theoretical limitation is their "black-box" nature. Connectionist models do not explicitly explain how they arrive at a conclusion; they provide a statistical probability without human readable reasoning. Furthermore, Representation Learning is heavily constrained by data dependency in such that the theory only works if the network is fed thousands of high quality, labeled examples to successfully learn the biological patterns.

2.2 Empirical Framework

The application of Artificial Intelligence for phytopathological diagnostics has grown exponentially. A review of existing literature reveals a strong progression from baseline feasibility studies to complex architectural comparisons, though a persistent gap in practical deployment remains.

Baseline Deep Learning Feasibility (Mohanty et al., 2016). Their objective was to determine if deep convolutional neural networks could reliably classify crop species and disease status from digital images. The researchers trained standard CNN architectures (AlexNet and GoogLeNet) on 54,306 images from the open source PlantVillage dataset, classifying 14 crop species and 26 diseases. The models achieved a groundbreaking accuracy of 99.35% on a held-out test set, demonstrating that deep learning could rival human expert diagnosis under controlled conditions. This study serves as the foundational proof of concept for the current research, validating that the Plant dataset is scientifically robust enough to train high performing diagnostic algorithms.

Architectural Optimization (Too et al., 2019). Their objective was to evaluate and compare the performance of various advanced deep learning architectures in identifying plant diseases to find the most computationally efficient model. The study utilized transfer learning to compare models including VGG16, Inception V4, ResNet and DenseNet on agricultural image datasets. The researchers found that DenseNet architectures consistently outperformed others achieving 99.75% accuracy while requiring significantly fewer computational parameters, as they effectively reused extracted features. This study directly informs the methodology of the current project, justifying the use of Transfer Learning and advanced parameter efficient architectures like ResNet or MobileNet, to maximize accuracy while minimizing computational load.

Mobile Deployment in African Agriculture (Ramcharan et al., 2019). The objective was to test the real world applicability of deep learning for detecting cassava diseases specifically within the East African agricultural context using low cost hardware. The researchers trained a MobileNet CNN architecture using Transfer Learning and deployed it onto an Android smartphone, testing its diagnostic capabilities on live cassava plants in the field. The mobile deployed model achieved an overall accuracy of 89%, proving that offline edge computing is viable for smallholder farming environments. This study bridges the gap between laboratory accuracy and field utility, strongly supporting the current study's goal of designing lightweight predictive models suitable for African agricultural frameworks.

Patterns, Contradictions and Gaps in the Findings: A clear pattern across the empirical literature is the undeniable statistical success of CNNs in classifying plant diseases, with laboratory accuracies consistently exceeding 95%. However, a contradiction arises when comparing laboratory performance to real-world deployment; models that score 99% on clean Plant data often experience performance drops as seen in (Ramcharan et al., 2019) 89% field accuracy when faced with varied lighting, backgrounds, and device constraints.

How the Current Project Builds on Previous Work:

The most glaring limitation in the current empirical literature is the lack of accessible deployment pipelines. The majority of studies end at algorithmic evaluation, leaving the predictive model locked in code repositories. The current research builds upon (Too et al., 2019). by utilizing their proven CNN optimization strategies, but it radically departs from standard theoretical studies by prioritizing software deployment. By wrapping the optimized CNN within an interactive Streamlit web dashboard, this project addresses the deployment gap, transitioning the research from a static accuracy metric into a functional, user accessible diagnostic tool for maize farmers.

2.3 Conceptual Framework

Independent Variable (Input Data):

The independent variable is the Maize Leaf Image Features. This includes the raw, unstructured RGB pixel data extracted from the maize leaf dataset. The variable is operationalized through specific biological markers present in the digital images, such as lesion geometry, color gradients and edge textures. These visual inputs act as the predictor variables that the algorithm will use to base its decisions.

Intervening Variables (Computational Processing):

The intervening variable represents the mechanism that explains the relationship between the independent and dependent variables. Here, it is the CNN Architecture and Hyperparameters. Guided by Representation Learning theory, this includes the convolution layers, pooling layers, fully connected layers, and activation functions. These computational layers act as the intervening bridge, autonomously extracting and mapping the low-level image features (IV) into high-level mathematical representations without human interference.

Dependent Variable (Predictive Outcome):

The dependent variable is the final outcome produced by the algorithmic pipeline, the Disease Classification. This variable is operationalized in two distinct outputs:


Categorical Label: The exact disease diagnosis e.g. Healthy, Common Rust, Northern Leaf Blight.

Confidence Score: A continuous percentage e.g., 97.5% representing the statistical probability of the model's accuracy.

Interconnections of the Variables

The relationship among these variables is strictly unidirectional and causal. The independent variable (raw image features) is fed into the intervening variable (the CNN processing engine). The intervening variable applies mathematical transformations (Transfer Learning and feature extraction) to interpret the visual data. This computational process directly dictates the value of the dependent variable which is the final predicted diagnosis and accuracy score.

Figure 2.1: Conceptual Framework for Maize Leaf Disease Detection Pipeline



CHAPTER 3: METHODOLOGY

3.1 Data Source and Description

Since deep learning algorithms require massive volumes of data to achieve high accuracy, this study utilizes secondary data. The primary data source is the PlantVillage dataset, a globally recognized, peer-reviewed, and open-source repository of agricultural imagery hosted on Kaggle. Specifically, the study will extract the subset of images dedicated to maize (*Zea mays*). This subset contains thousands of high-resolution RGB photographs of maize leaves, categorically labeled by plant pathologists into four distinct classes:

1. Healthy
2. Common Rust (*Puccinia sorghi*)
3. Northern Leaf Blight (*Exserohilum turcicum*)
4. Cercospora Leaf Spot (*Cercospora zeae-maydis*)

3.2 Data Preprocessing and Augmentation

Raw digital images cannot be directly fed into a neural network; they must first be converted into a standardized mathematical format. Using Python's TensorFlow and Keras libraries, the data pipeline will involve the following preprocessing steps:

1. Resizing: All images will be programmatically resized to a uniform dimension like 224x224 pixels to match the input requirements of the selected neural network.
2. Normalization: The pixel intensity values, which originally range from 0 to 255, will be mathematically scaled down to a range between 0 and 1. This ensures computational efficiency and speeds up the algorithm's convergence during training.
3. Data Augmentation: To prevent the model from "memorizing" the images (overfitting), dynamic data augmentation will be applied. This involves writing code to randomly rotate, flip, zoom, and shift the images during training, forcing the algorithm to learn the underlying disease patterns rather than the specific orientation of the leaf.

3.3 Model Architecture and Training Strategy

The core computational engine of this study is a Convolutional Neural Network (CNN). However, building and training a CNN from scratch is computationally expensive. Therefore, this study will employ Transfer Learning.

1. **The Architecture:** The study will utilize a pre-trained architecture, such as MobileNetV2 or ResNet50. These models have already been trained on millions of generic images and possess highly optimized convolutional filters for detecting edges, shapes, and textures.
2. **Fine Tuning:** The base layers of the model will be frozen, and a new, fully connected classification head will be added. This new head will be trained exclusively on the maize leaf dataset.
3. **Training Parameters:** The dataset will be split into three sets: Training (80%), Validation (10%), and Testing (10%). The model will be compiled using the Adam optimizer and a Categorical Cross-Entropy loss function, iterating over multiple epochs until the validation loss stabilizes.

3.4 Performance Evaluation Metrics

The model's predictive power will not be evaluated on the data it learned from, but strictly on the 10% hidden "Testing" dataset. Performance will be quantified using standard data science metrics via the scikit-learn library.

Accuracy: The overall percentage of correctly diagnosed leaf images.

Precision and Recall: Precision will measure how many of the predicted disease cases were actually correct, while Recall will measure how many actual disease cases the model successfully found.

Confusion Matrix: A visual grid will be generated to map the exact false positives and false negatives for example checking if the model frequently confuses Rust with Blight.

3.5 Model Deployment

Following successful evaluation, the trained model's computational weights will be exported as a standalone executable file like model.h5. To bridge the gap between code and practical utility, a graphical user interface (GUI) will be engineered using the Streamlit Python library. This web-based dashboard will allow users to upload a new, unclassified image of a maize leaf, pass the image through the saved algorithm in the backend, and instantly display the predicted disease class alongside a statistical confidence score.

3.6 Ethical Considerations

Because this study relies exclusively on open-source, de-identified secondary data (digital images of plants), it does not involve human subjects, animal testing, or personal data collection. Consequently, the research requires no ethical clearance regarding participant confidentiality or physical harm. Furthermore, to uphold the highest standards of academic integrity, all datasets, pre-trained algorithmic architectures, and Python libraries will be properly referenced and cited.

REFERENCES:

- Khan, F., Zafar, N., Tahir, M. N., Aqib, M., Waheed, H., & Haroon, Z. (2023). A mobile-based system for maize plant leaf disease detection and classification using deep learning. *Frontiers in Plant Science*, 14(May), 1–18.
<https://doi.org/10.3389/fpls.2023.1079366>
- Mahlein, A.-K. (2016). Present and Future Trends in Plant Disease Detection. *Plant Disease*, 100(2), 1–11. <https://doi.org/10.1007/s13398-014-0173-7.2>
- Medler, D. a. (1998). A Brief History of Connectionism. *Neural Computing Surveys*, 1(2), 18–72. <https://doi.org/10.1007/BF01889762>
- Mohanty, S. P., Hughes, D. P., & Salathé, M. (2016). Using deep learning for image-based plant disease detection. *Frontiers in Plant Science*, 7(September), 1–10.
<https://doi.org/10.3389/fpls.2016.01419>
- NBS, [National Bureau of Statistics], & OCGS, [Office of the Chief Government Statistician]. (2024). *Tanzania Integrated Labour Force Survey 2024 Report*.
- Noroozi, M., Pirsiavash, H., & Favaro, P. (2017). Representation Learning by Learning to Count. *Proceedings of the IEEE International Conference on Computer Vision, 2017-October*, 5899–5907. <https://doi.org/10.1109/ICCV.2017.628>
- Nsibo, D. L., Barnes, I., & Berger, D. K. (2024). Recent advances in the population biology and management of maize foliar fungal pathogens *Exserohilum turcicum*, *Cercospora zeina* and *Bipolaris maydis* in Africa. *Frontiers in Plant Science*, 15(August), 1–23.
<https://doi.org/10.3389/fpls.2024.1404483>
- Ramcharan, A., McCloskey, P., Baranowski, K., Mbilinyi, N., Mrisho, L., Ndalawha, M., Legg, J., & Hughes, D. P. (2019). A mobile-based deep learning model for cassava disease diagnosis. *Frontiers in Plant Science*, 10(March), 1–8.
<https://doi.org/10.3389/fpls.2019.00272>
- Savary, S., Willocquet, L., Pethybridge, S. J., Esker, P., McRoberts, N., & Nelson, A. (2019). The global burden of pathogens and pests on major food crops. *Nature Ecology & Evolution*, 3(3), 430–439. <https://doi.org/10.1038/s41559-018-0793-y>
- Too, E. C., Yujian, L., Njuki, S., & Yingchun, L. (2019). A comparative study of fine-tuning deep learning models for plant disease identification. *Computers and Electronics in Agriculture*, 161(October 2017), 272–279.
<https://doi.org/10.1016/j.compag.2018.03.032>
- Wu, J. (2017). Introduction to Convolutional Neural Networks. *Introduction to Convolutional Neural Networks*, 1–31.
https://web.archive.org/web/20180928011532/https://cs.nju.edu.cn/wujx/teaching/15_CNN.pdf